
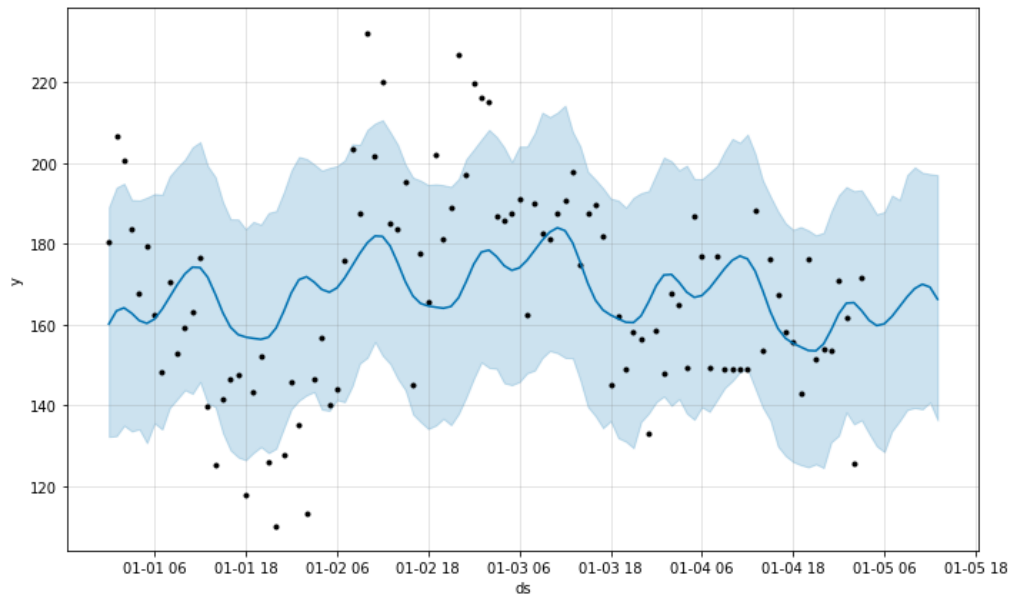


Multistep prediction



Related information

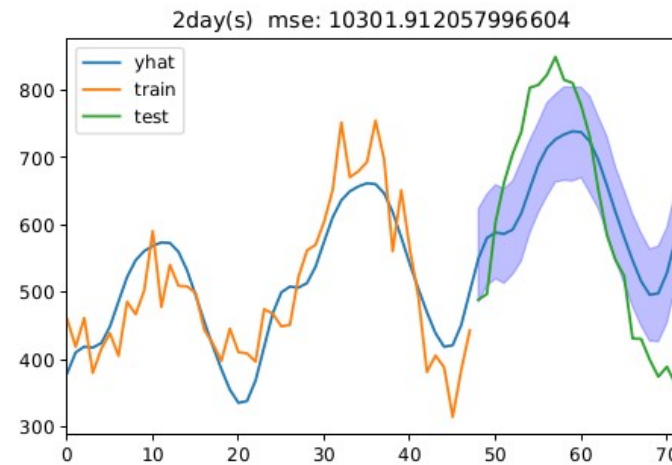
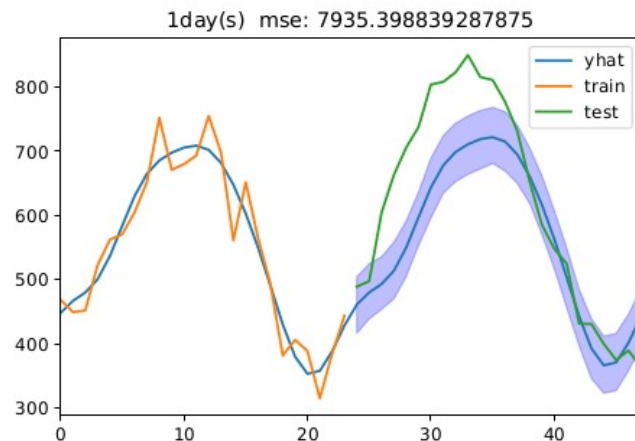
- Time series usually have yearly seasonality, weekly seasonality and daily seasonality
- $y_{\hat{}} = trend + [daily_seasonality] + [weekly_seasonality] + [yearly_seasonality]$
- Example:



- Tool: Prophet https://facebook.github.io/prophet/docs/quick_start.html

Overview of my ideas

- N period points before current point C must have a trend, using this trend to predict M period points later ($\text{trend}(N) \rightarrow \text{valueOf}(M)$)
- Set $\text{lower}N$ and $\text{upper}N$, then let N differs in $[\text{lower}N, \text{upper}N]$, different trends can be get, so we get multi $\text{valueOf}(M)$ sequences.
- Give the $\text{valueOf}(M)$ weights , make weighted summation, finally got the $\text{hat_valueOf}(M)$



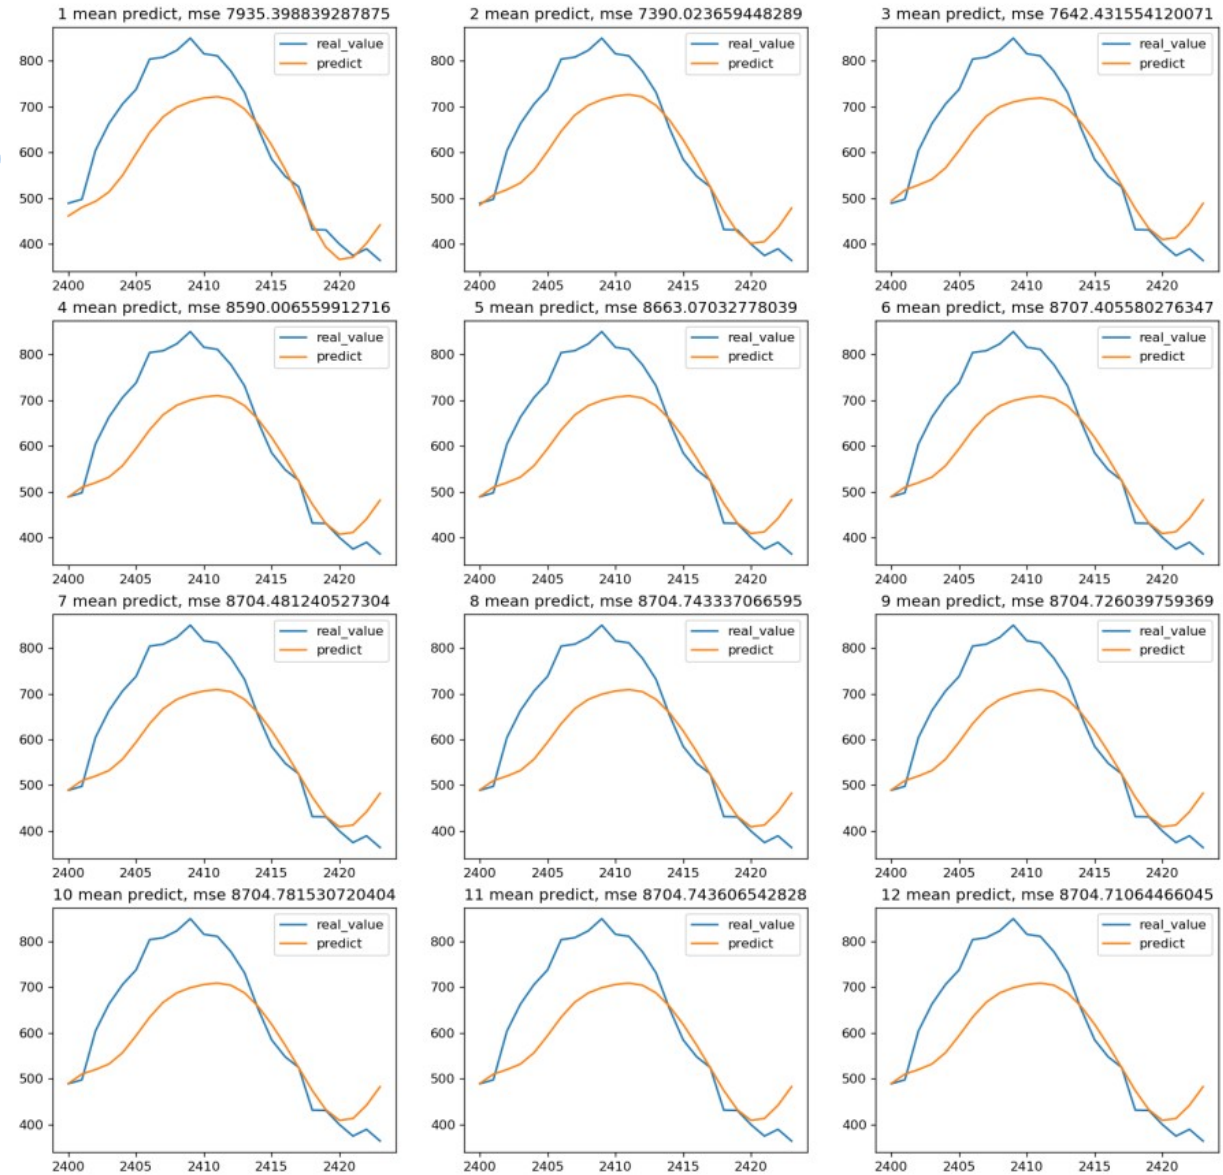
• More>>

Pseudocode

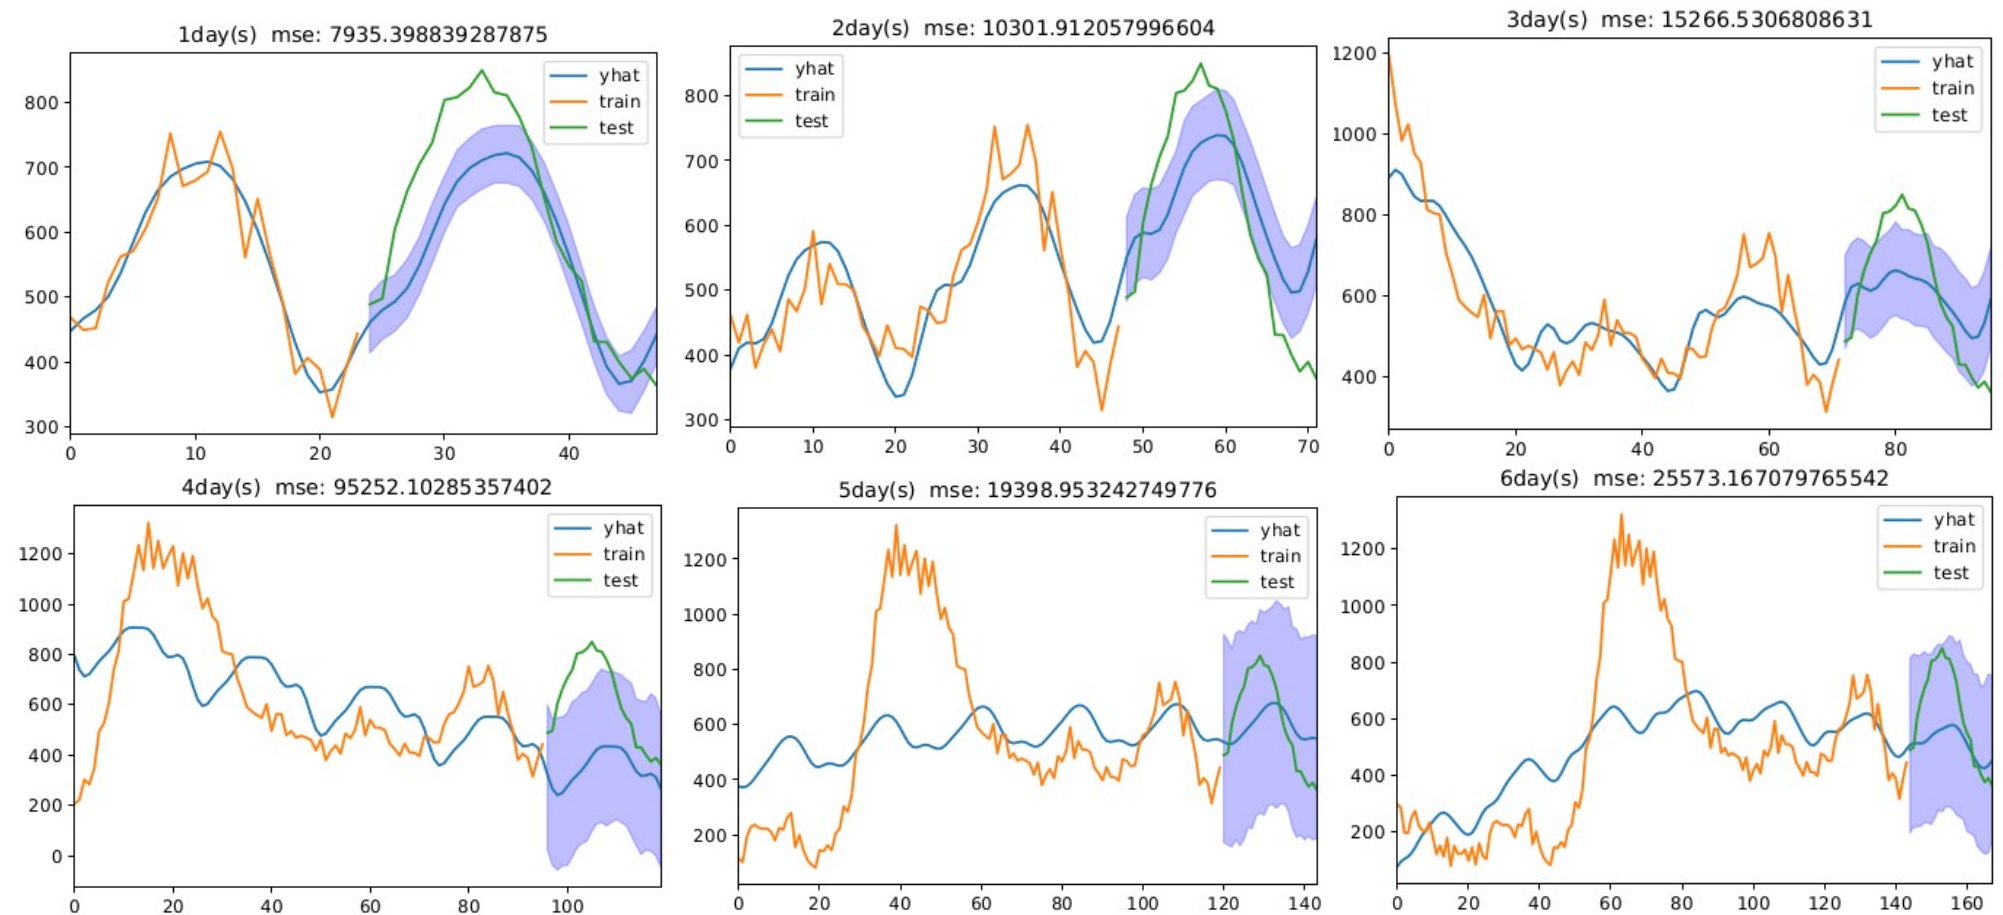
- **For** N in range($lowerN$, $upperN$, step):
 - ValueOfM_pre_list = []
 - Trend = get_trend(N _points_sequence)
 - ValueOfM_pre = next_M_points_predict(Trend, C , M)
 - ValueOfM_pre_list.append(ValueOfM_pre)
- **For** $weight$ in weights, ValueOfM_pre in ValueOfM_pre_list:
 - $ValueOfM_hat += weight * ValueOfM_pre$
- **Finally** get the predict of next M periods values $ValueOfM_hat$

Result obtained 1

- $C = 2015-04-11\ 00:00:00$
- N in range($[24, 24 \cdot 12 + 1, 24]$)
- $M = 24$



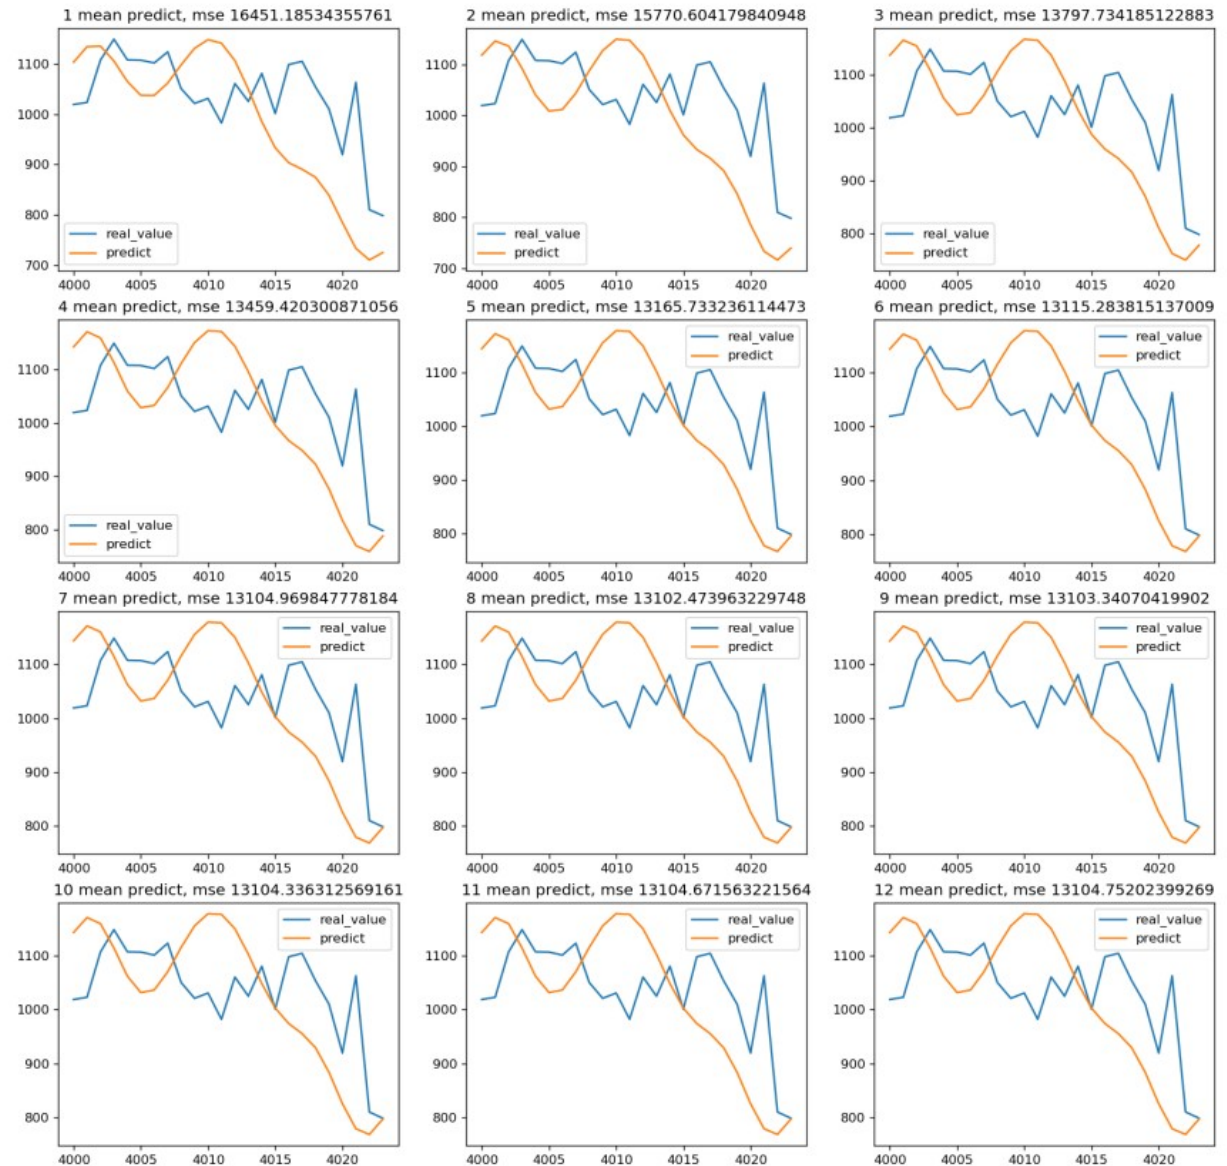
Result obtained 1 details



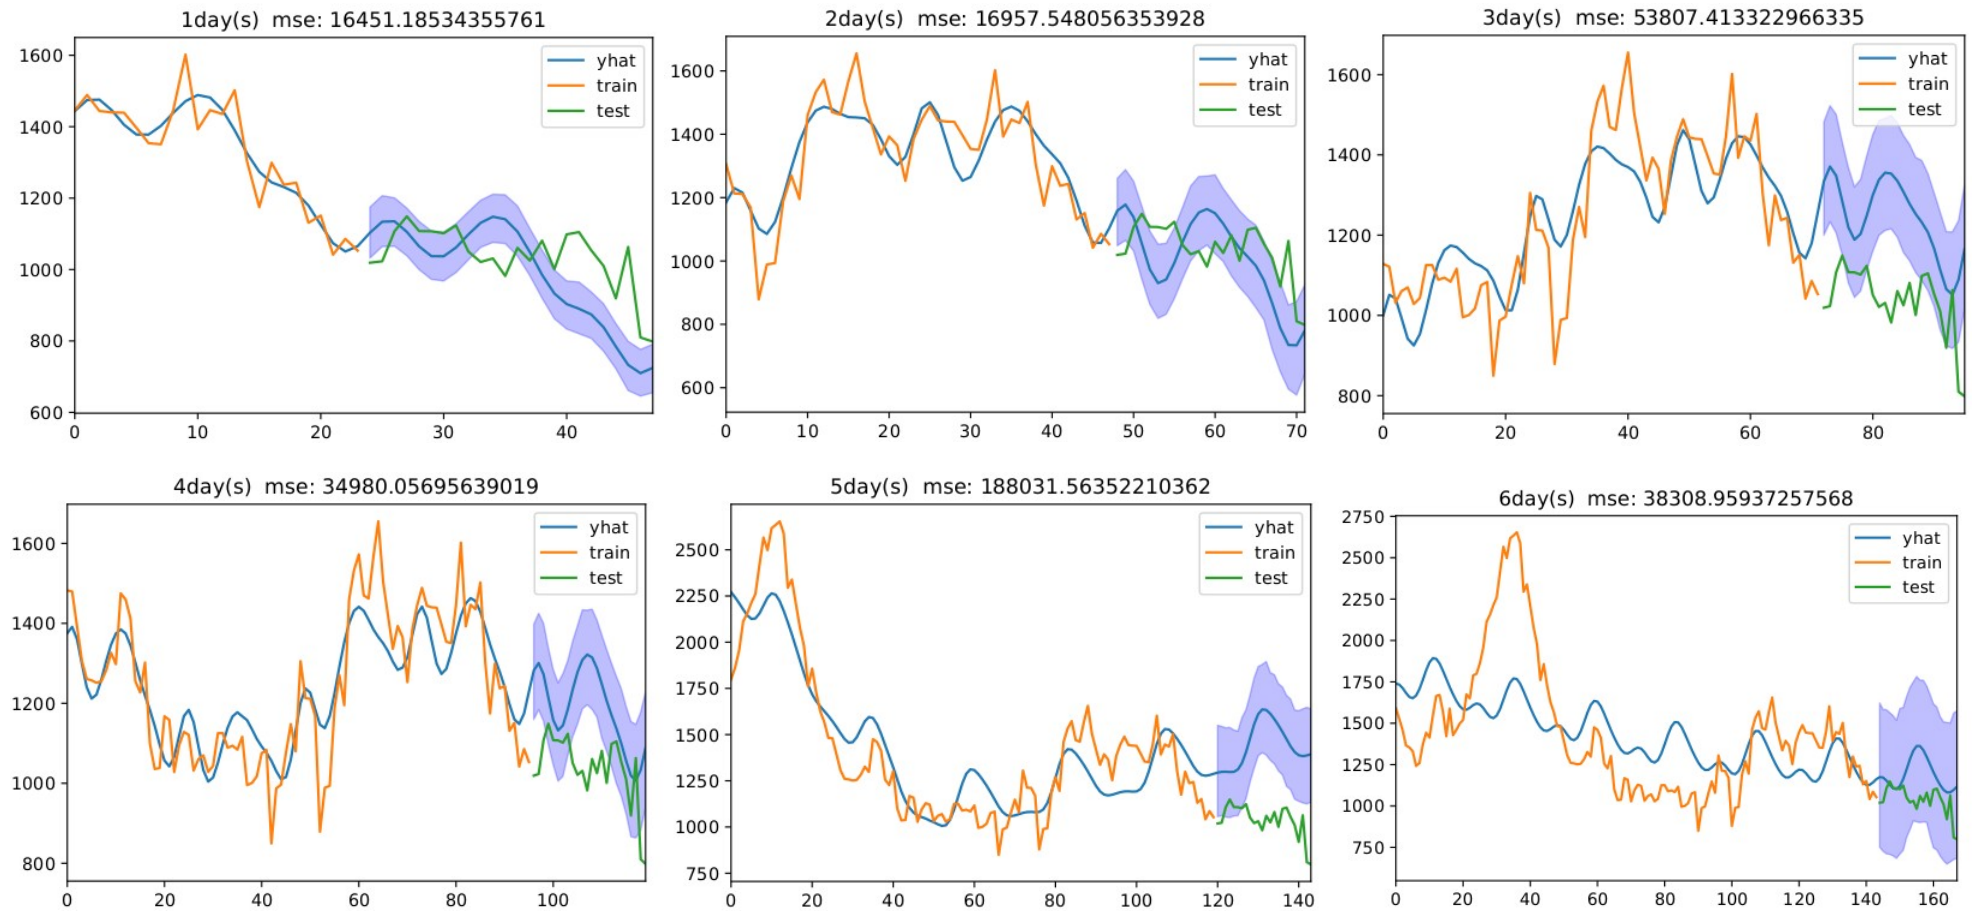
.....

Result obtained 2

- **C** = 2015-06-16 16:00:00
- **N** in range([24,24*12+1,24])
- **M** = 24



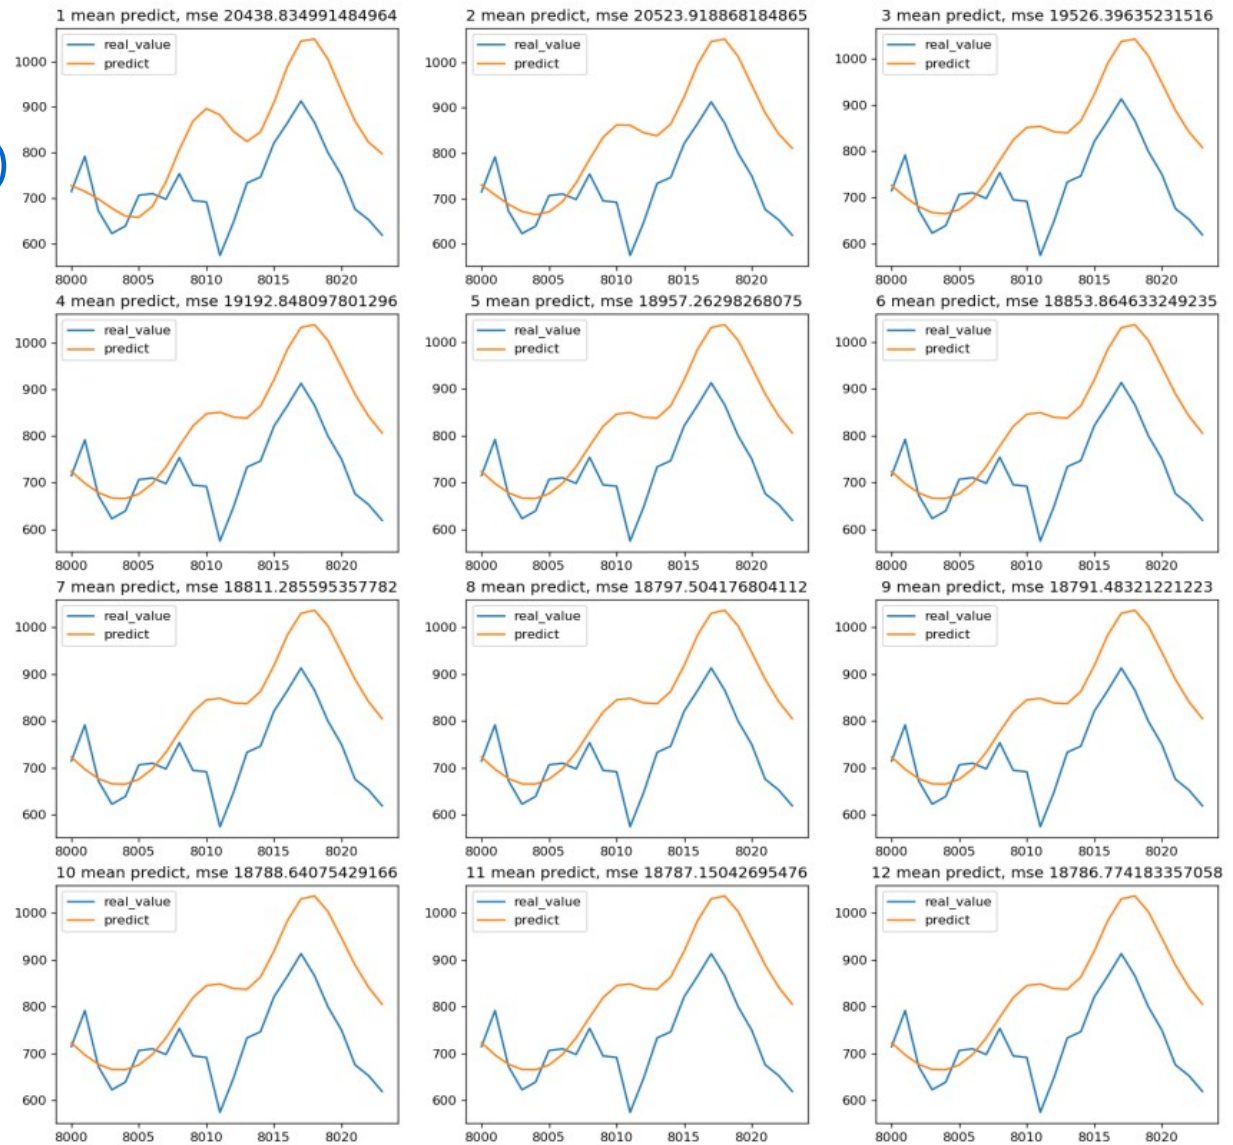
Result obtained 2 details




.....

Result obtained 3

- $C = 2015-11-30\ 08:00:00$
- N in range($[24, 24 \times 12 + 1, 24]$)
- $M = 24$



Result obtained 3 details



.....

Weight strategies

- Let n = Number of next period points prediction sequences
 - Strategy1: share the same weight $1/n$
 - Strategy2: $e^i / \sum(e^j)$ (i, j in range(n))
- Code:

```
In [6]: def get_weights(n,weight_type='same'):
        if weight_type is 'same':
            return np.array(n*[1.0/n])
        elif weight_type is 'softmax':
            denominator = np.sum(np.exp([i for i in range(n)]))
            # print(denominator)
            numerator = np.exp(np.abs(np.sort([-i for i in range(n)])))
            # print(numerator)
            return numerator / denominator
        else:
            pass
        print('same weight\n',get_weights(12,'same'),'\\n')
        print('softmax weight\n',get_weights(12,'softmax'))
```

```
same weight
[0.08333333 0.08333333 0.08333333 0.08333333 0.08333333 0.08333333
 0.08333333 0.08333333 0.08333333 0.08333333 0.08333333 0.08333333]
```

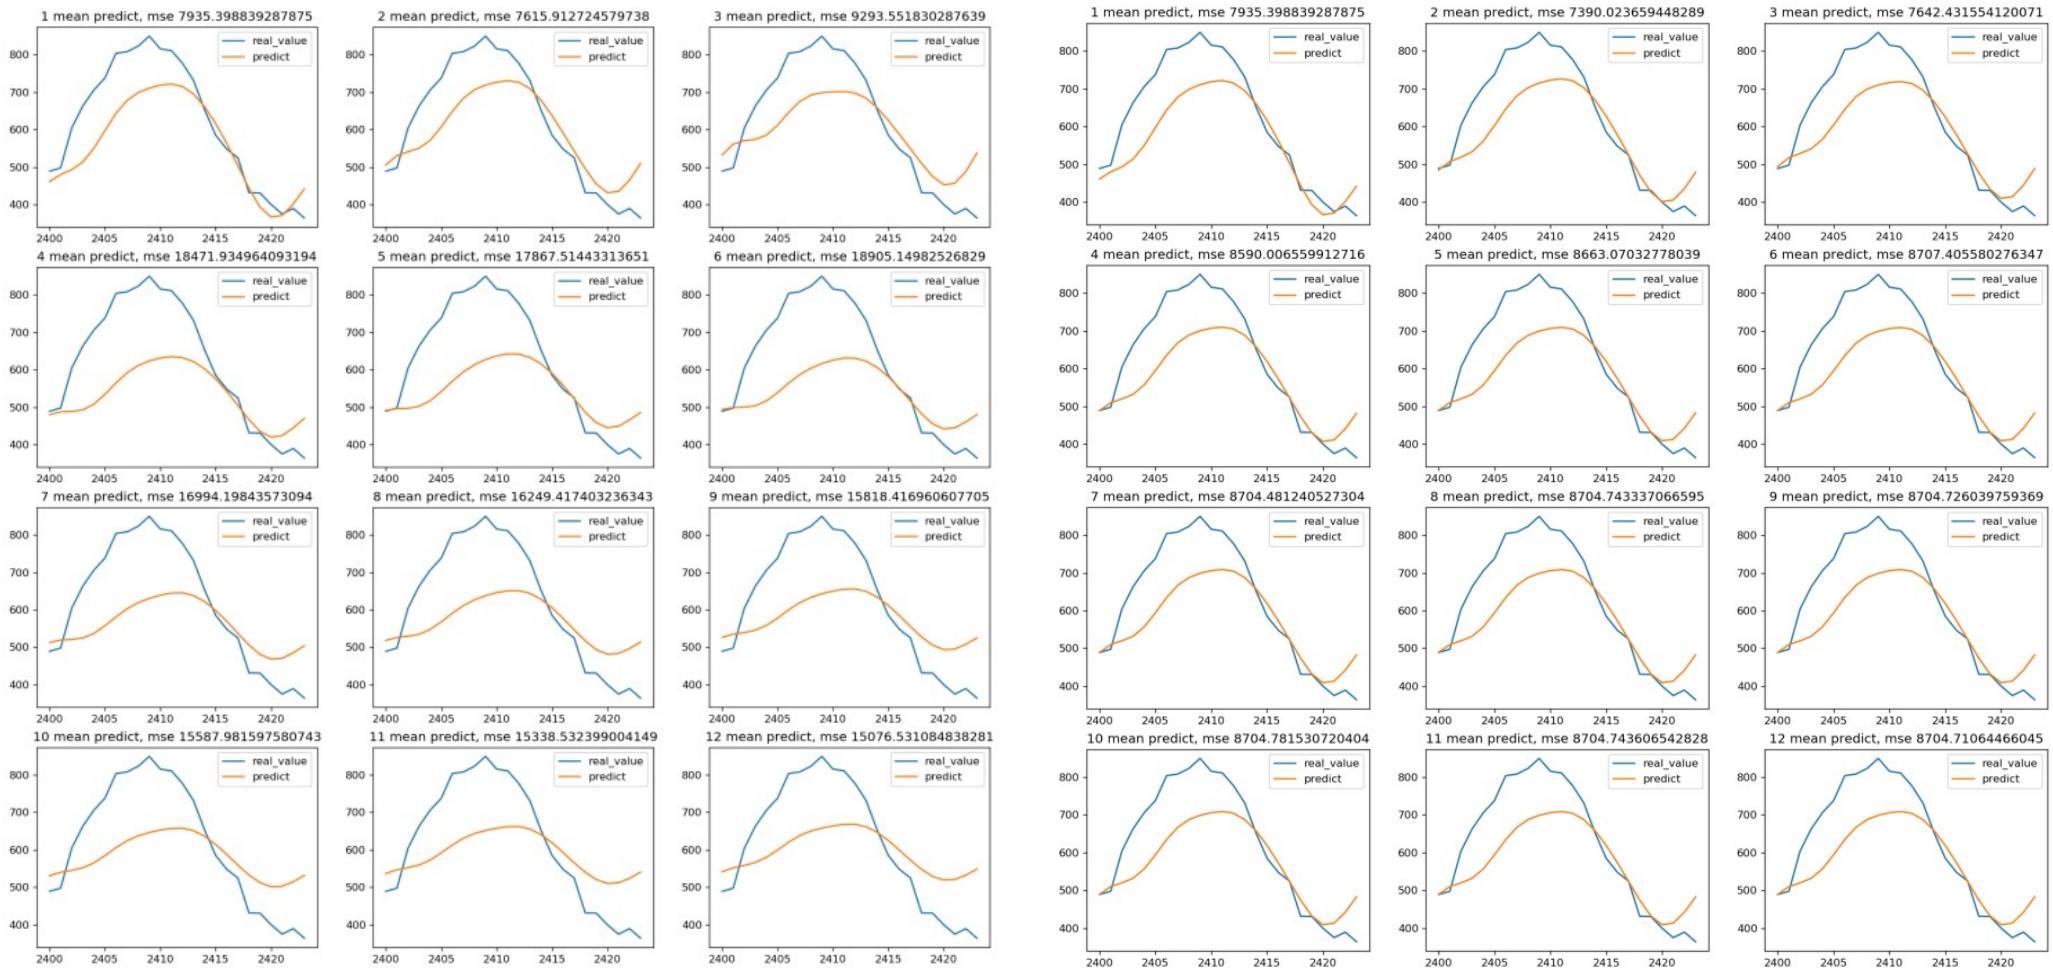
```
softmax weight
[6.32124443e-01 2.32545587e-01 8.55487405e-02 3.14716228e-02
 1.15777630e-02 4.25922099e-03 1.56687984e-03 5.76422879e-04
 2.12054127e-04 7.80103536e-05 2.86984053e-05 1.05575533e-05]
```

Comparison of two strategies

‘same’

VS

‘softmax’



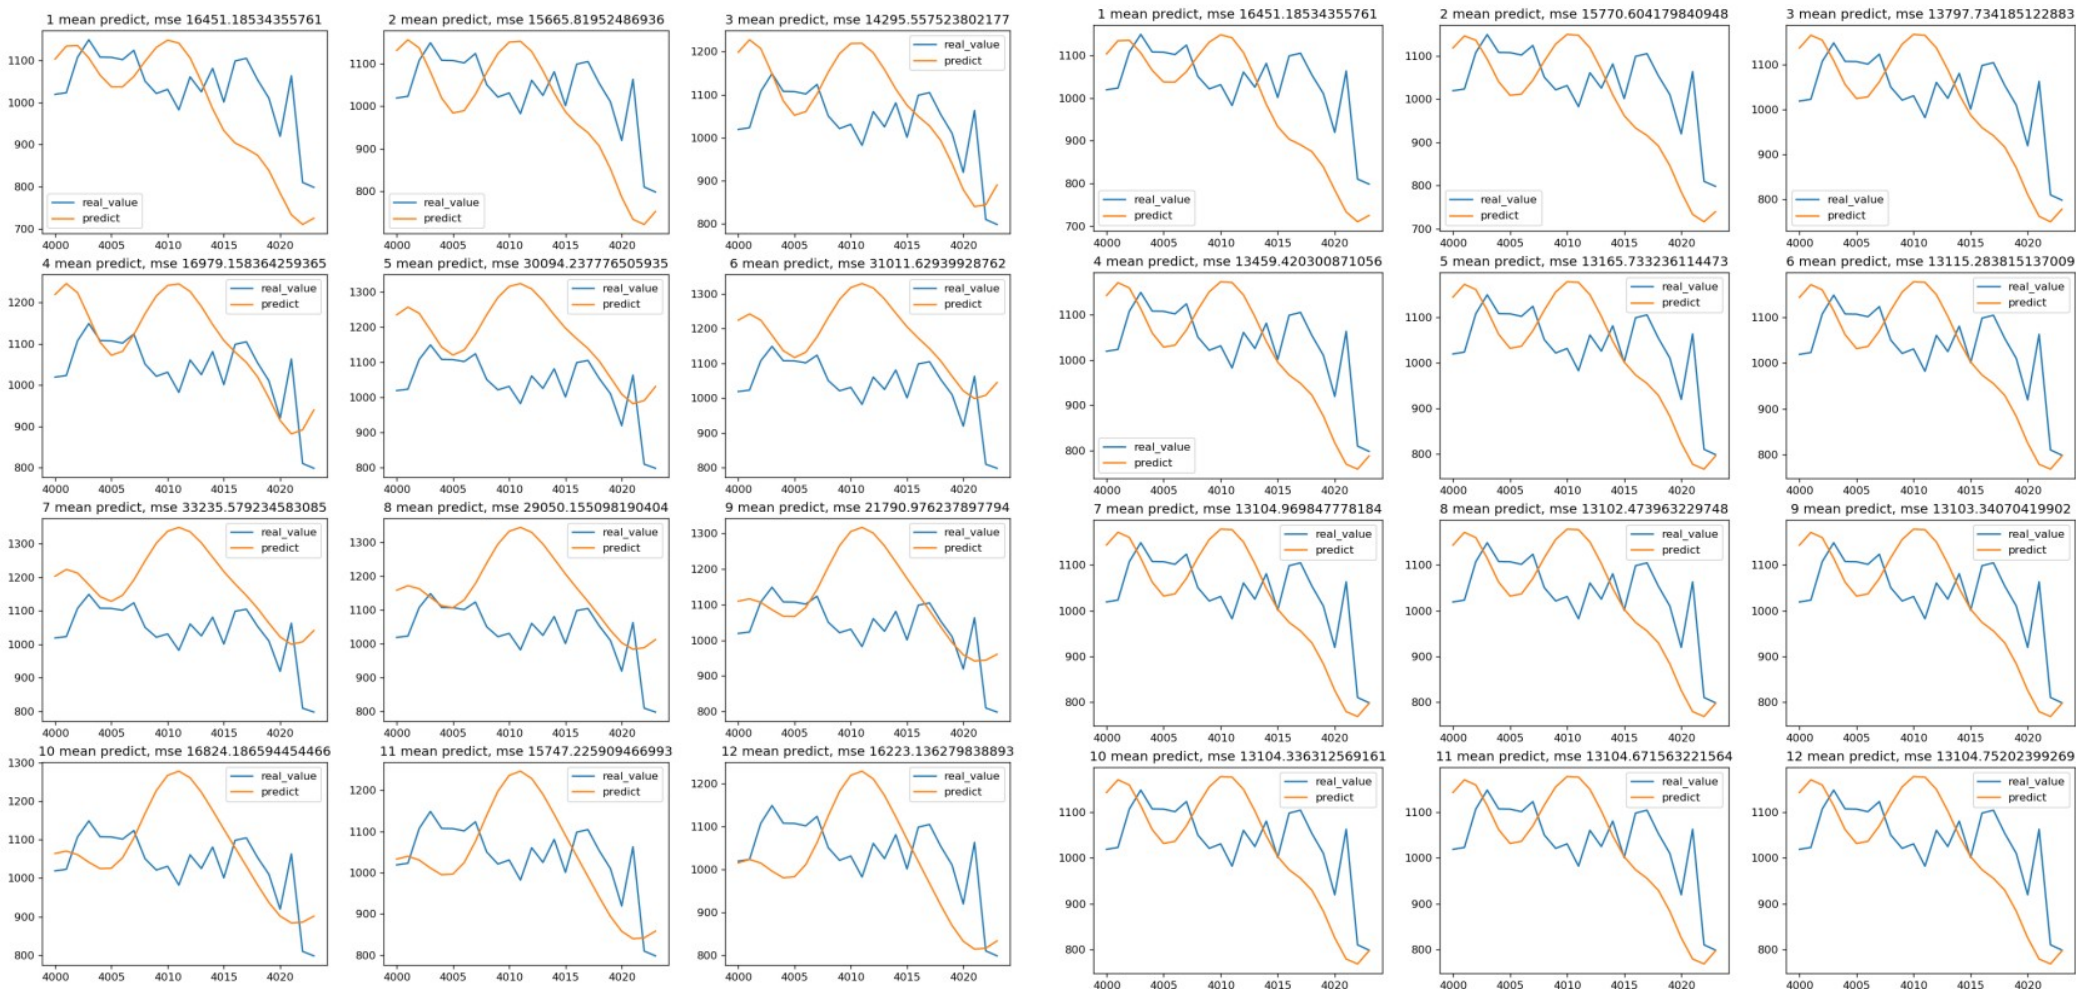
Comparison of two strategies

-

‘same’

VS

‘softmax’

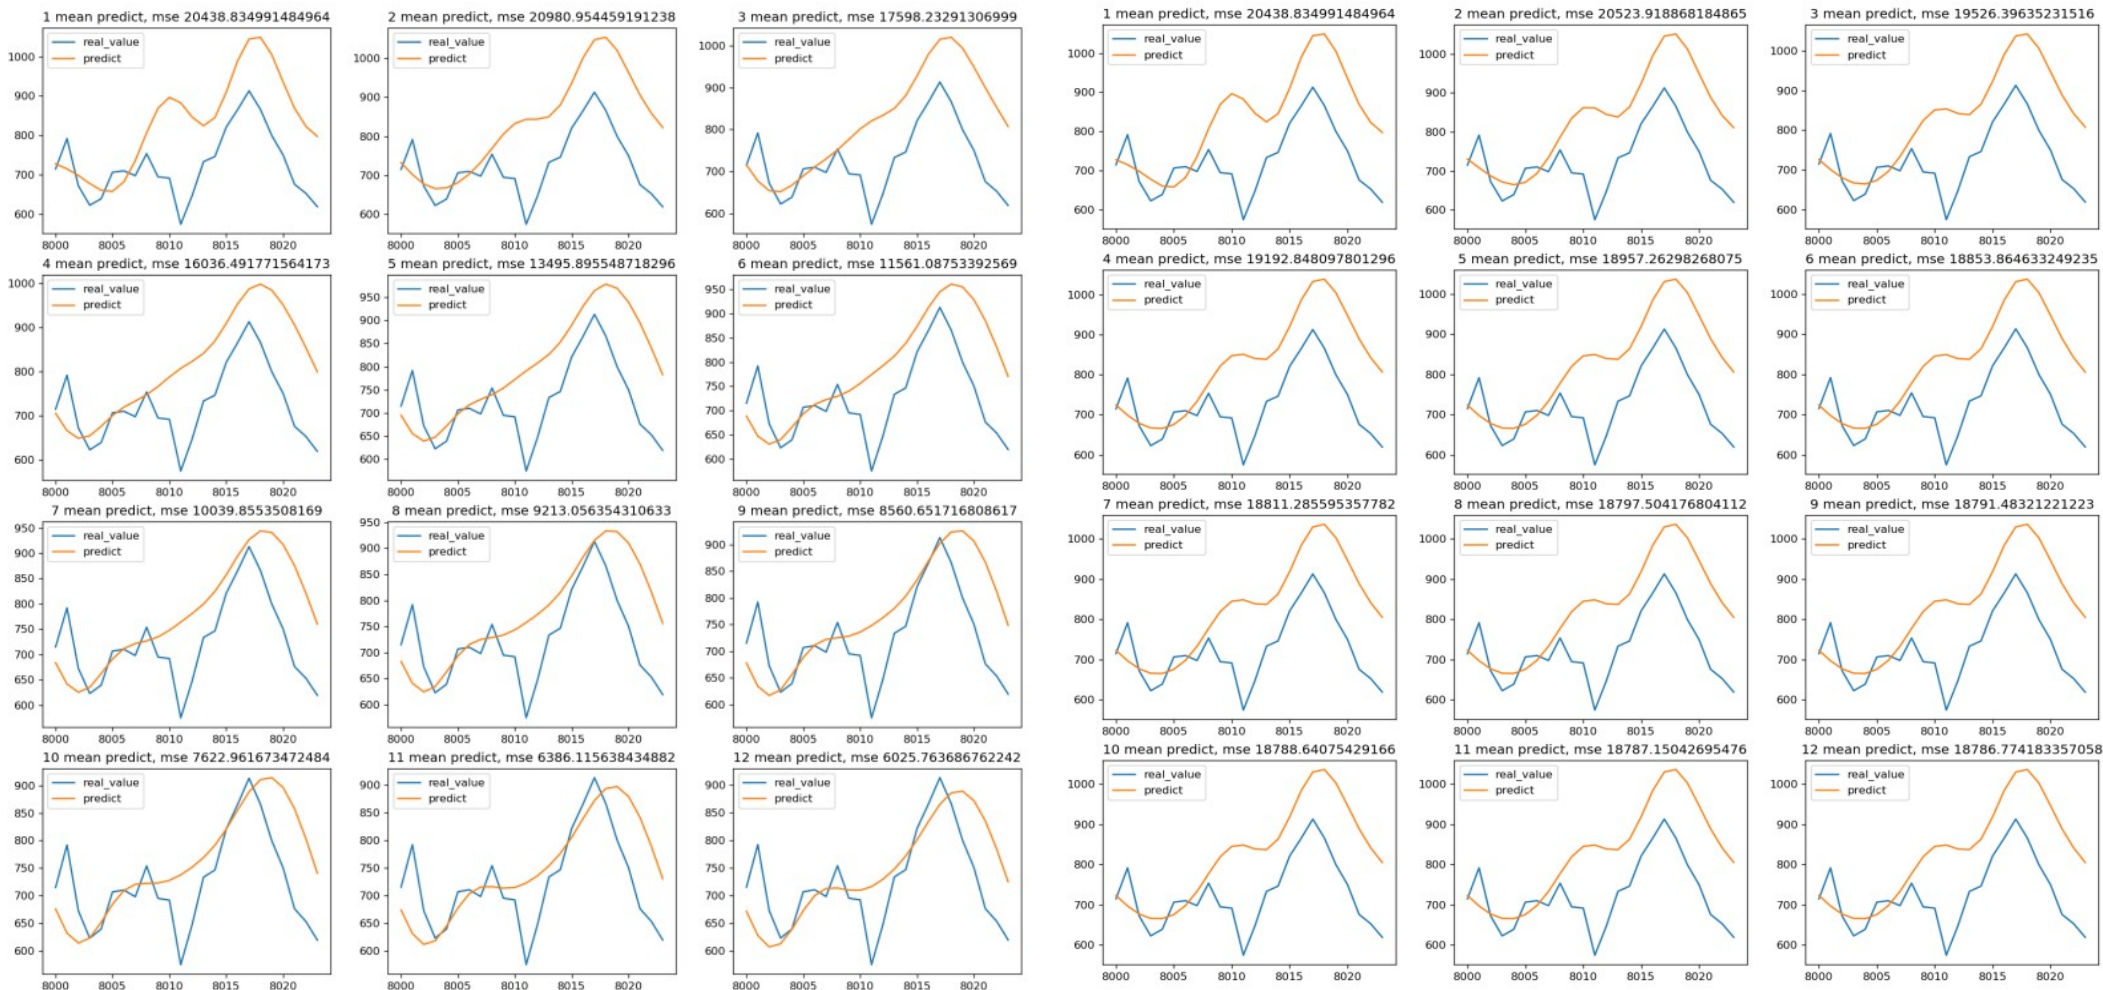


Comparison of two strategies

‘same’

VS

‘softmax’

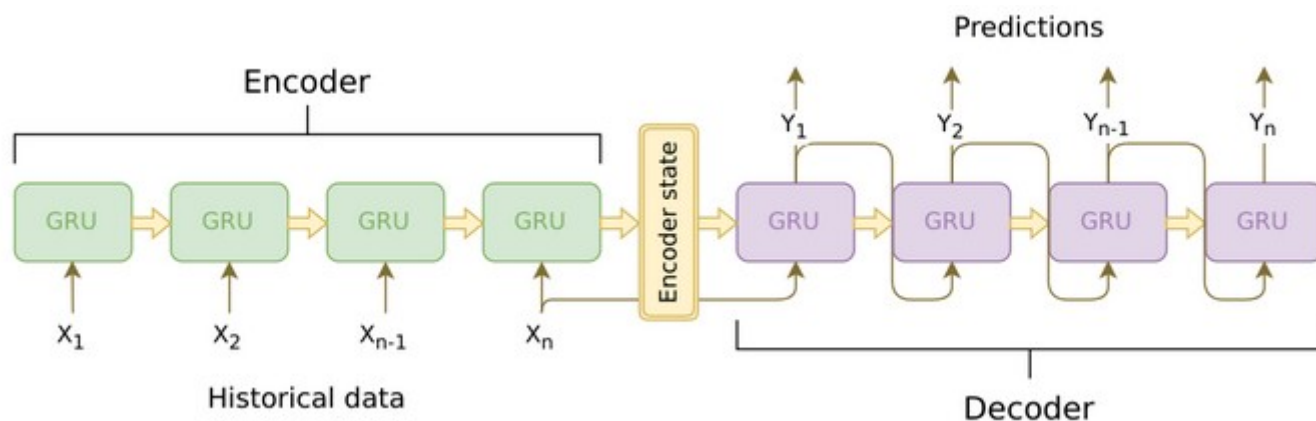


Advantages and disadvantages

- Advantages:
 - The model is very simple, just focus on the data of the recent Npoints, it only takes very little training data and time.
 - Models can 'remember' the impact of data points far away from themselves on current data trends
 - The effects of recent and distant trends on current data trends can be adjusted with different weights
- Disadvantages:
 - Finding the optimal weight assignment strategy may be difficult
 - The trend of different Npoints sequence may be quite different, which is unfavorable for the prediction of the results.

Solution assumption

- Improve weights assignment strategy
- Construct a supervised learning model to 'learn' weights
- Using Encoder, Decoder(seq2seq)
 - Mapping a fixed length sequence to a consecutive fixed length sequence(with GRU or LSTM layers)



- Demo implementation
 - <https://www.kaggle.com/c/web-traffic-time-series-forecasting/discussion/43795#latest-567361>

Improved weight strategy

- Lower pass filter





图 1 理想低通滤波器冲激响应图




Improved weight strategy

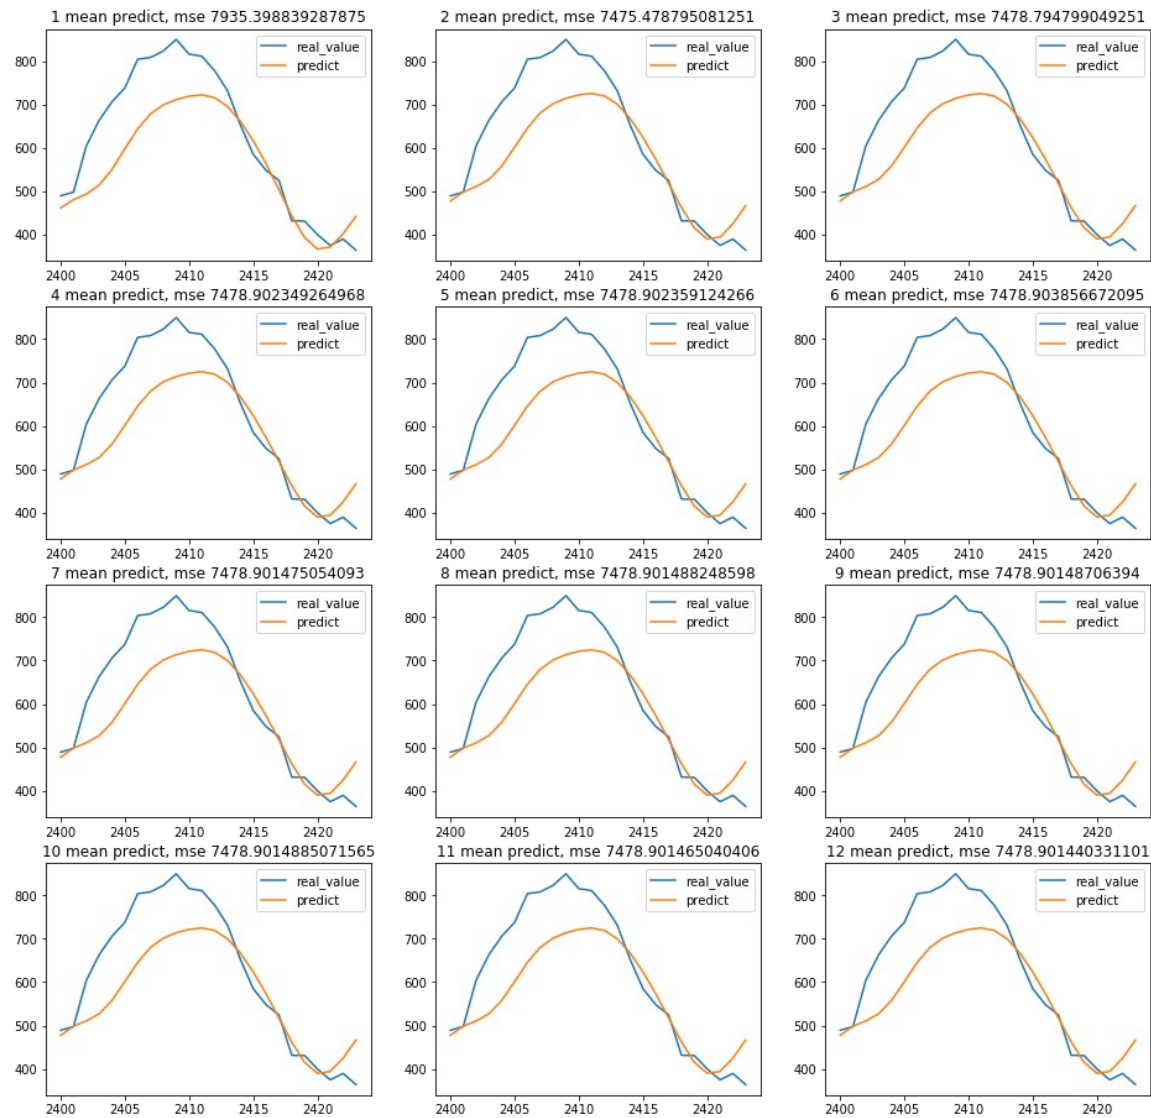
- Lower pass filter

```
In [155]: def lower_pass_filter(n,period,amplitude,shift):
    res = []
    for i in range(int(n / period)+1):
        weight_decay = (1 / (10 ** 0.5)) ** i
        if i == 0:
            res.append(shift + weight_decay * amplitude * np.cos([i*(np.pi/2)/period for i in range(period)]))
        else:
            res.append(shift + weight_decay * amplitude * np.cos([-np.pi/2 + i*(np.pi)/period for i in range(period)])
    return res
res = np.array(lower_pass_filter(12,4,10,0.2)).reshape(-1)
plt.plot(res)
denominator = np.sum(np.exp(res))
numerator = np.exp(res)
res = numerator / denominator
res
```

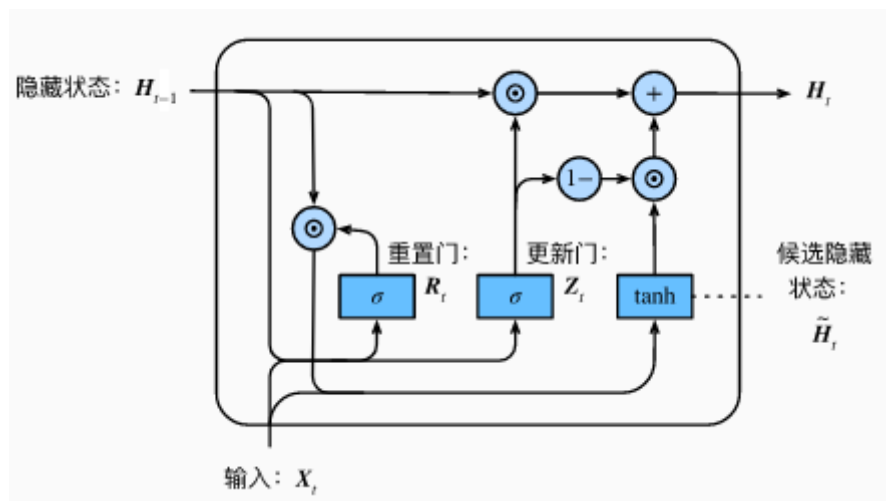
```
Out[155]: array([6.55658680e-01, 3.06260386e-01, 3.50476359e-02, 1.36680332e-03,
 2.97668580e-05, 2.78512685e-04, 7.03222462e-04, 2.78512685e-04,
 2.97668580e-05, 6.03706108e-05, 8.09147093e-05, 6.03706108e-05,
 2.97668580e-05, 3.72258133e-05, 4.08384237e-05, 3.72258133e-05])
```



Improved results



GRU(gated recurrent unit)



- Reset Gate & Update Gate

$$R_t = \sigma(X_t W_{xr} + H_{t-1} W_{hr} + b_r),$$

$$Z_t = \sigma(X_t W_{xz} + H_{t-1} W_{hz} + b_z),$$

- Candidate hidden state

$$\tilde{H}_t = \tanh(X_t W_{xh} + (R_t \odot H_{t-1}) W_{hh} + b_h),$$

- Hidden state

$$H_t = Z_t \odot H_{t-1} + (1 - Z_t) \odot \tilde{H}_t.$$

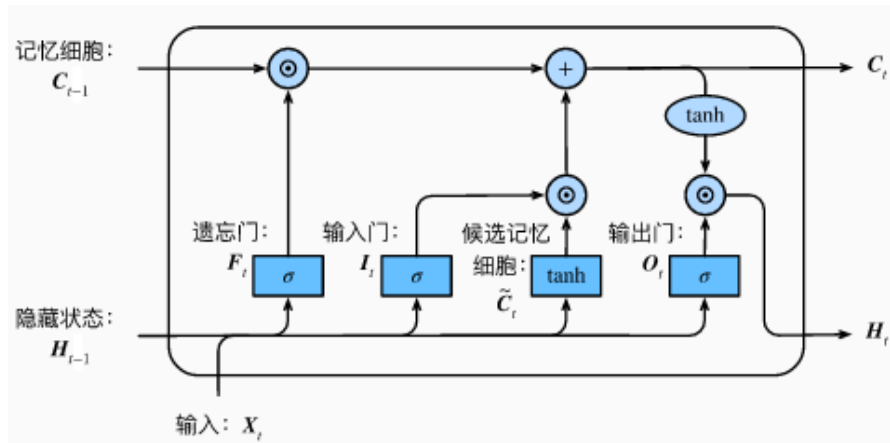
- Summary

- Reset gate can be used to discard some historical information unrelated to prediction
- Reset gate helps capture short-term dependencies in the time series
- Update gate helps capture long-term dependencies in time series

- References

- <https://arxiv.org/pdf/1409.1259.pdf>
- <https://arxiv.org/pdf/1412.3555.pdf>

LSTM(long short-term memory)



- Input, Forget, Output Gate

$$\begin{aligned} I_t &= \sigma(X_t W_{xi} + H_{t-1} W_{hi} + b_i), \\ F_t &= \sigma(X_t W_{xf} + H_{t-1} W_{hf} + b_f), \\ O_t &= \sigma(X_t W_{xo} + H_{t-1} W_{ho} + b_o), \end{aligned}$$

- Candidate Memory Cell

$$\tilde{C}_t = \tanh(X_t W_{xc} + H_{t-1} W_{hc} + b_c),$$

- Memory Cell

$$C_t = F_t \odot C_{t-1} + I_t \odot \tilde{C}_t.$$

- Hidden state

$$H_t = O_t \odot \tanh(C_t).$$


- Summary

- The hidden layer output of LSTM includes H_t and C_t . Only H_t is passed to the output layer
- LSTM network can cope with the gradient attenuation problem in the cyclic neural network and better capture the dependence of the time step distance in the time series.

- References

- <https://link.springer.com/content/pdf/10.1007%2F978-3-642-24797-2.pdf>

Encoder&Decoder(seq2seq)



- Encoder

$$h_t = f(x_t, h_{t-1})$$

$$c = q(h_1, \dots, h_T).$$

- Decoder

$$s_{t'} = g(y_{t'-1}, c, s_{t'-1}).$$

- **How to choose c ????**

- Summary


- Encoder transforms the hidden state of each time step into a background variable through a custom function q
- Decoder $P(y_{t'} \mid y_1, \dots, y_{t'-1}, c)$

- References

- <https://arxiv.org/pdf/1406.1078.pdf>
- <https://arxiv.org/pdf/1409.3215.pdf>

Attention mechanism in seq2seq

- The attention should be assigned differently by different steps.
 - Global Attention



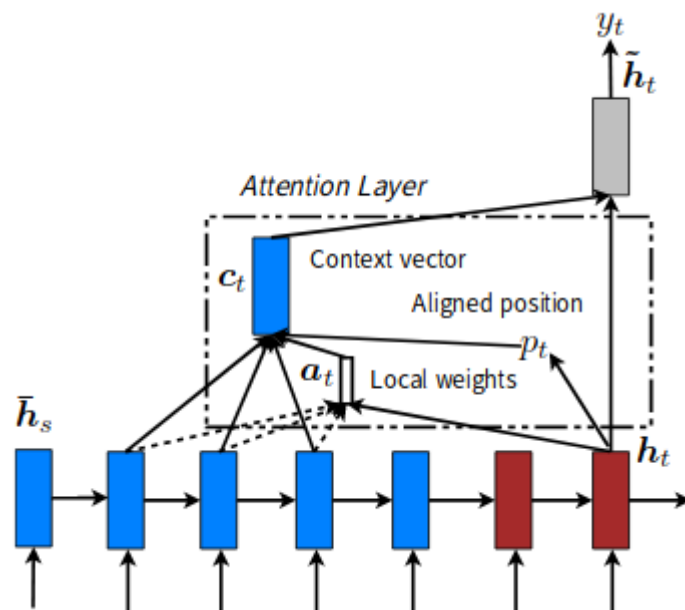
$$\begin{aligned} a_t(s) &= \text{align}(h_t, \bar{h}_s) \\ &= \frac{\exp(\text{score}(h_t, \bar{h}_s))}{\sum_{s'} \exp(\text{score}(h_t, \bar{h}_{s'}))} \end{aligned}$$

$$\text{score}(h_t, \bar{h}_s) = \begin{cases} h_t^\top \bar{h}_s & \text{dot} \\ h_t^\top W_a \bar{h}_s & \text{general} \\ v_a^\top \tanh(W_a [h_t; \bar{h}_s]) & \text{concat} \end{cases}$$

Figure 2: **Global attentional model** – at each time step t , the model infers a *variable-length* alignment weight vector a_t based on the current target state h_t and all source states \bar{h}_s . A global context vector c_t is then computed as the weighted average, according to a_t , over all the source states.

Attention mechanism in seq2seq

- The attention should be assigned differently by different steps.
 - Local Attention



$$p_t = S \cdot \text{sigmoid}(\mathbf{v}_p^\top \tanh(\mathbf{W}_p \mathbf{h}_t)), \quad (9)$$

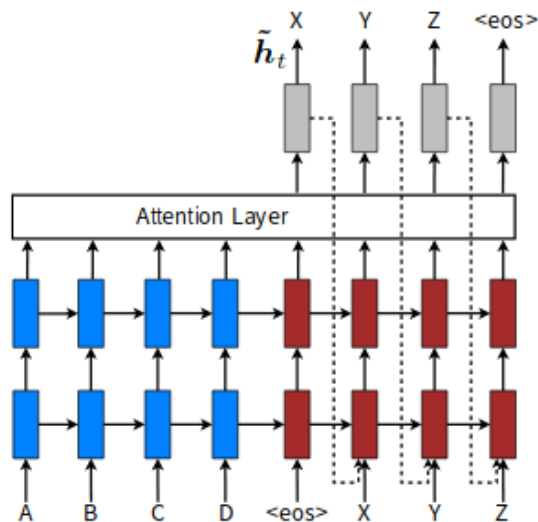
\mathbf{W}_p and \mathbf{v}_p are the model parameters which will be learned to predict positions. S is the source sentence length. As a result of sigmoid, $p_t \in [0, S]$. To favor alignment points near p_t , we place a Gaussian distribution centered around p_t . Specifically, our alignment weights are now defined as:

$$\mathbf{a}_t(s) = \text{align}(\mathbf{h}_t, \bar{\mathbf{h}}_s) \exp\left(-\frac{(s - p_t)^2}{2\sigma^2}\right) \quad (10)$$

We use the same align function as in Eq. (7) and the standard deviation is empirically set as $\sigma = \frac{D}{2}$. Note that p_t is a real number; whereas s is an integer within the window centered at p_t .¹⁰

Attention mechanism in seq2seq

- The attention should be assigned differently by different steps.
 - Input-feeding Approach



- References
 - <https://arxiv.org/pdf/1508.04025.pdf>

Figure 4: **Input-feeding approach** – Attentional vectors \tilde{h}_t are fed as inputs to the next time steps to inform the model about past alignment decisions.

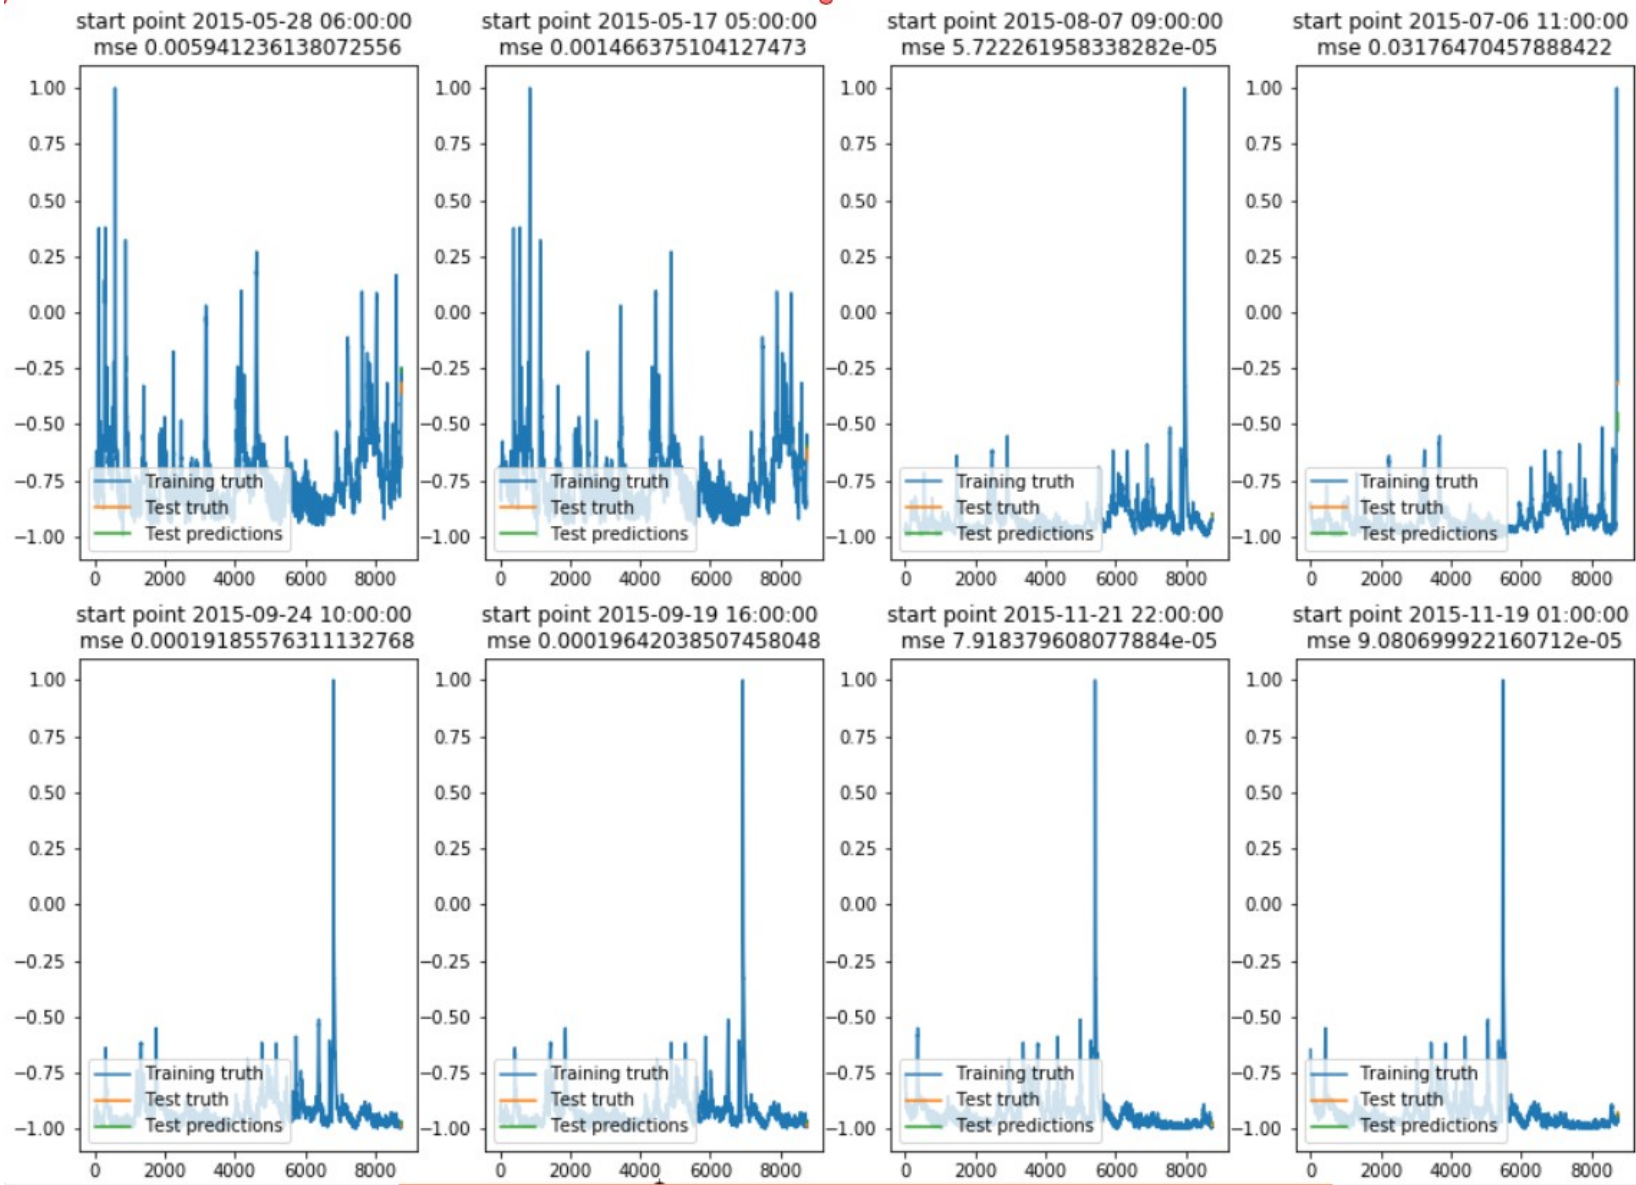
Using seq2seq

- Data preprocessing: MinMaxScaler((-1,1))
- Make data into: sequences → sequences

	current_before_3	current_before_2	current_before_1	current	current_next_1	current_next_2
3	-0.964450	-0.961909	-0.96248	-0.964150	-0.965671	-0.964560
4	-0.961909	-0.962480	-0.96415	-0.965671	-0.964560	-0.966197

- We can build data set seq(len: N) → seq(len: M) in any $N \& M$
- Have a try with $N = 3$, $M = 3$

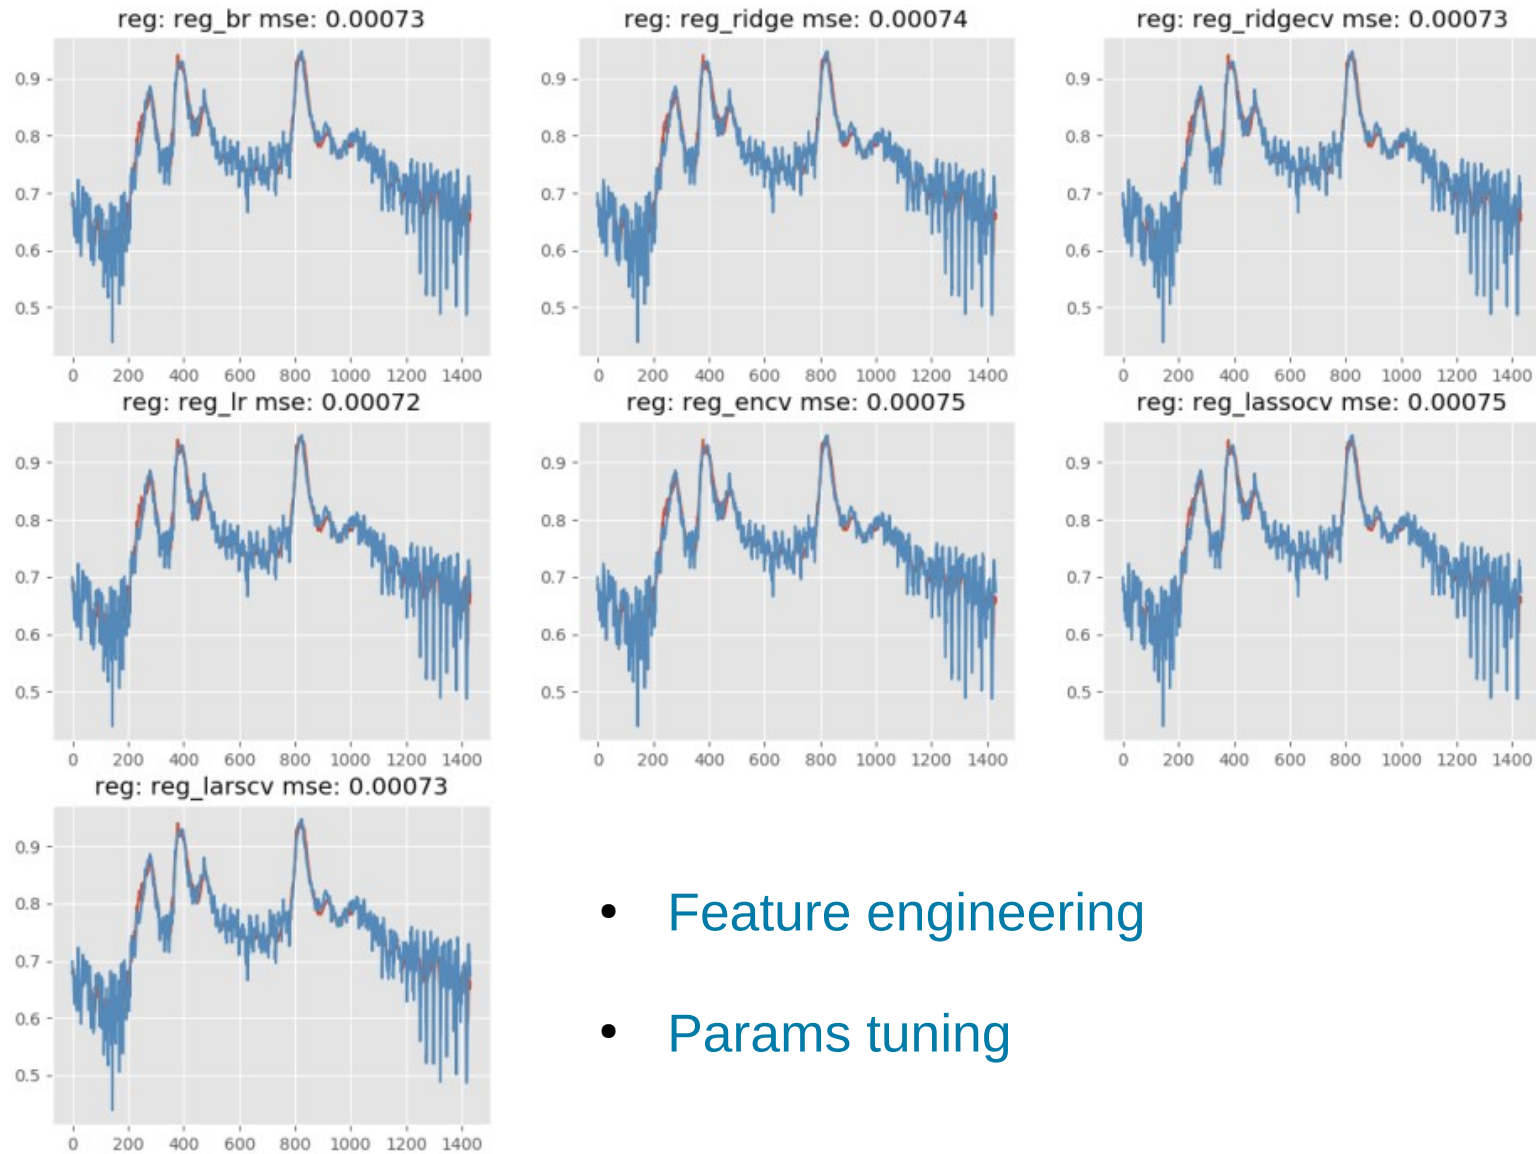
Result obtained



Single step prediction

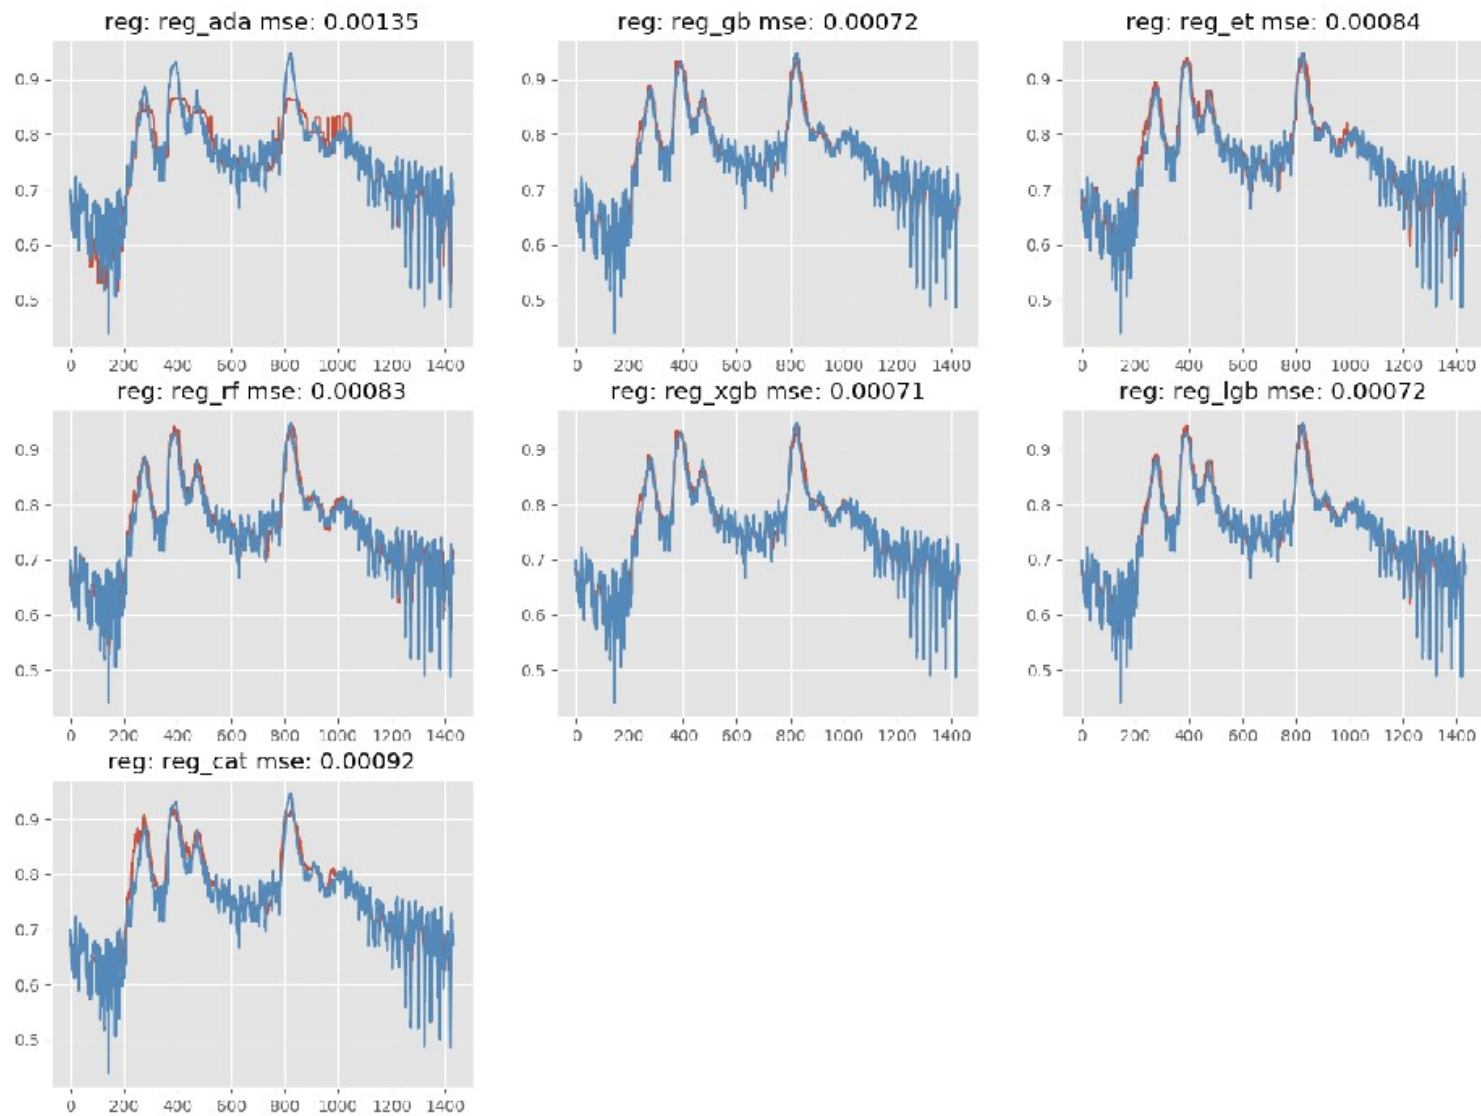


Linear model



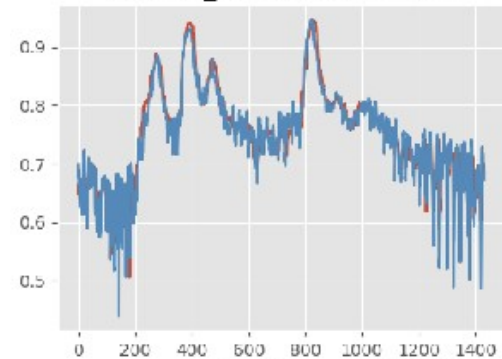
- Feature engineering
- Params tuning

Tree Based model




Bagging


reg: reg_bag1 mse: 0.00079




reg: reg_bag2 mse: 0.00112




reg: reg_bag3 mse: 0.00069




reg: reg_bag4 mse: 0.00076




reg: reg_bag5 mse: 0.00073




reg: reg_bag6 mse: 0.0007



reg: reg_bag7 mse: 0.00067




reg: reg_bag8 mse: 0.00094




```
reg_bag1 = BaggingRegressor()  
reg_bag2 = BaggingRegressor(base_estimator=reg_ada)  
reg_bag3 = BaggingRegressor(base_estimator=reg_gb)  
reg_bag4 = BaggingRegressor(base_estimator=reg_et)  
reg_bag5 = BaggingRegressor(base_estimator=reg_rf)  
reg_bag6 = BaggingRegressor(base_estimator=reg_xgb)  
reg_bag7 = BaggingRegressor(base_estimator=reg_lgb)  
reg_bag8 = BaggingRegressor(base_estimator=reg_cat)
```


Voting


reg: reg_vote1 mse: 0.00073



reg: reg_vote2 mse: 0.0007



reg: reg_vote3 mse: 0.00071



```
linear_models = []
linear_models.append(('reg_br', reg_br))
linear_models.append(('reg_ridge', reg_ridge))
linear_models.append(('reg_ridgecv', reg_ridgecv))
linear_models.append(('reg_lr', reg_lr))
linear_models.append(('reg_encv', reg_encv))
linear_models.append(('reg_lassocv', reg_lassocv))
linear_models.append(('reg_larscv', reg_larscv))
reg_vote1 = VotingRegressor(estimators=linear_models)
```

```
tree_based_models = []
tree_based_models.append(('reg_ada', reg_ada))
tree_based_models.append(('reg_gb', reg_gb))
tree_based_models.append(('reg_et', reg_et))
tree_based_models.append(('reg_rf', reg_rf))
tree_based_models.append(('reg_xgb', reg_xgb))
tree_based_models.append(('reg_lgb', reg_lgb))
tree_based_models.append(('reg_cat', reg_cat))
reg_vote2 = VotingRegressor(estimators=tree_based_models)
```

```
bagging_models = []
bagging_models.append(('reg_bag1', reg_bag1))
bagging_models.append(('reg_bag2', reg_bag2))
bagging_models.append(('reg_bag3', reg_bag3))
bagging_models.append(('reg_bag4', reg_bag4))
bagging_models.append(('reg_bag5', reg_bag5))
bagging_models.append(('reg_bag6', reg_bag6))
bagging_models.append(('reg_bag7', reg_bag7))
bagging_models.append(('reg_bag8', reg_bag8))
reg_vote3 = VotingRegressor(estimators=bagging_models)
```